

Lessons learned from deploying Rust apps in production

Leon Oosterwijk

Akuna Capital

Oct 2024

Overview

- Introduction
- Use-case
- Outcome
- Implementation details
- Lessons learned
- Things I wish for

Introduction

- Akuna Capital
 - HFT Trading firm
 - HQ in Chicago
 - Trading APAC markets from Sydney
 - Focused on Index Option Market making
 - Participant on NSE, the world's largest exchange by volume
 - Large technology footprint in C++
- Leon Oosterwijk
 - Tech lead for 'Core Systems'
 - Experience across telco, ecommerce, banking and trading
 - Mostly focused on 'business' languages (Java, Python)
 - 2 years experience with Rust -
 - 1 in production

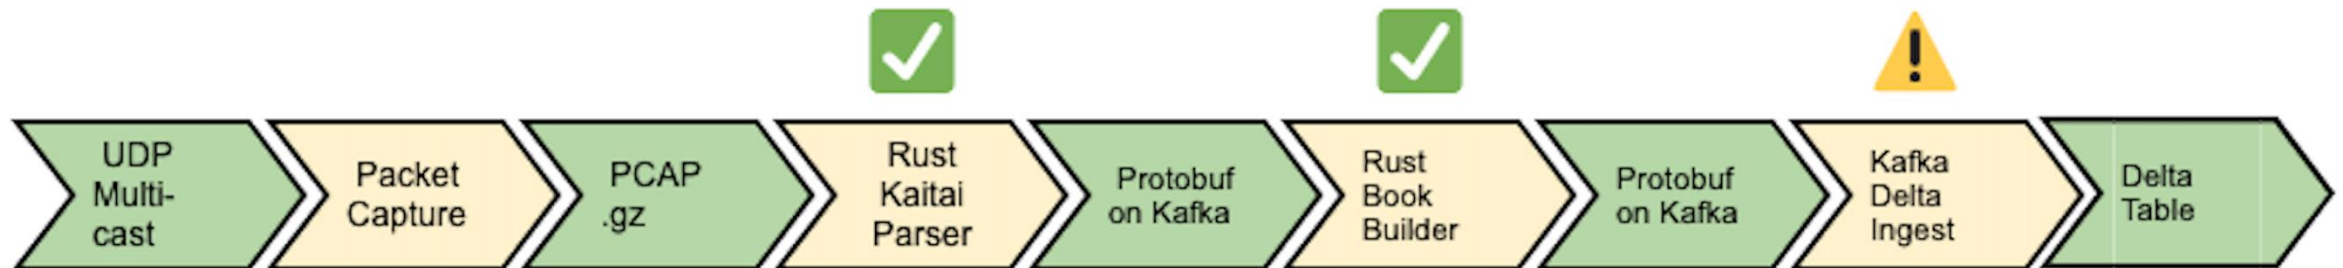
Use Case

- Data pipeline for NSE Market data
 - Approx 7-10 Billion events per day (1TB compressed)
- Existing pipeline was brittle, slow and expensive
- Book building requires a lot of state



Outcome

- Deployed PCAP Parser
- Deployed Book Builder
- Kafka Delta Ingest in dev – pending prod deploy
- Resource usage approx 1/3 of previous solution (parser)
- No more delay from realtime



Implementation Details

- Libraries

- [Etherparse](#), [Flatbuffers](#), [Kaitai](#) for pcap parsing
- [Rust-rdkafka](#) for kafka
- [Protofish](#) and [protobuf](#)
- Tokio and rayon for async / parallelism
- Grafana for observability ([metrics-rs](#))

- Deployment

- Container in EKS (Exploring graviton deployment)
- Cron and Lag based scaling

Lessons learned

- Many rust libraries are dormant or immature
- `async` is infectious
- If it compiles, it runs
- Controlling memory footprint with `librdkafka` is tricky
- For long running apps, leaking to ``static` is sometimes best
- `.clone()` first, optimize later
- Only store references in structs if you really really have to

Things I wish for. . .

- Flink style as-of join capability
- Actor-like framework
- Faster compiles (or is it the linking that's slow?)
- Reflection